

CS-GY 6923: Lecture 14

Diffusion and Transformers

NYU Tandon School of Engineering, Akbar Rafiey

- Final Dec 19. Same time and location.
- Homework 5 due on Dec 10.

Diffusion Model

Auto-encoder/VAE, GAN approach: Input noise, map directly to image and vice versa.

Diffusion: Slowly move from noise to image and vice versa.

Denoising Diffusion Probabilistic Models

Jonathan Ho
UC Berkeley
jonathanho@berkeley.edu

Ajay Jain
UC Berkeley
ajayj@berkeley.edu


Pieter Abbeel
UC Berkeley
pabbeel@cs.berkeley.edu

Abstract

We present high quality image synthesis results using diffusion probabilistic models, a class of latent variable models inspired by considerations from nonequilibrium thermodynamics. Our best results are obtained by training on a weighted variational bound designed according to a novel connection between diffusion probabilistic

How diffusion models work

- Forward Process:
 - Gradually add noise to data until it becomes pure noise.
- Reverse Process:
 - Train a neural network to remove the noise step by step.




Key Question: How do we predict and reverse noise effectively?

Mathematical formulation

Forward Process (Adding Noise):

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$$



- β_t : Noise schedule.
- After T steps, for large enough T , $\mathbf{x}_T$ is pure noise.

Cumulative Noise:


$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$$

with retention factor $\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$.

Mathematical formulation

Reverse Process (Denoising):

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \Sigma_{\theta}(\mathbf{x}_t, t))$$



- μ_{θ} : Predicted mean of the clean image.
- Σ_{θ} : Predicted variance (optional).

Training objective:

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{\mathbf{x}_0, t, \epsilon} [\|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|^2]$$

Diffusion models vs VAEs

- Recall the goal of VAE was to have a probabilistic representation of latent space attributes. This was done in one shot, from image to Gaussian distributions.
- VAE decoder does the reverse in one shot. Takes Gaussian noise and returns an image.
- Diffusion model doing more or less the same but with many careful intermediate noising and denoising steps.
- There are fundamental differences ...

Diffusion process

- A diffusion process is a stochastic Markov process having continuous path
- stochastic Markov process: future state will only depend on the current state, knowing the past does not change anything
- continuous: no jumps
- Allows to transition from complex distributions to simple distributions

Forward process: a closer look

Forward process: $q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1-\beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$

Why does it converge to a simple distribution ?

Suppose the process is $\mathbf{x}_t = \alpha\mathbf{x}_{t-1} + \beta\mathcal{N}(0, \mathbf{I})$.

What values of α and β makes sense for the process?

- $\alpha = 0, \beta = 1$?
- $\alpha \geq 1, \beta \leq 1$?
- Seems something like $\alpha = \sqrt{0.99}, \beta = \sqrt{0.01}$ is appropriate.

Forward process: a closer look

Let's check if the following α and β converges:

$$\mathbf{x}_t = \sqrt{1 - \beta} \mathbf{x}_{t-1} + \sqrt{\beta} \mathcal{N}(0, I)$$

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta} \mathbf{x}_{t-1} + \sqrt{\beta} \mathcal{N}(0, I) \\ &= \sqrt{1 - \beta} (\sqrt{1 - \beta} \mathbf{x}_{t-2} + \sqrt{\beta} \mathcal{N}(0, I)) + \sqrt{\beta} \mathcal{N}(0, I) \\ &= (\sqrt{1 - \beta})^2 \mathbf{x}_{t-2} + \sqrt{1 - \beta} \sqrt{\beta} \mathcal{N}(0, I) + \sqrt{\beta} \mathcal{N}(0, I)\end{aligned}$$

Forward process: a closer look

Let's check if the following α and β converges:

$$\mathbf{x}_t = \sqrt{1 - \beta} \mathbf{x}_{t-1} + \sqrt{\beta} \mathcal{N}(0, I)$$

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta} \mathbf{x}_{t-1} + \sqrt{\beta} \mathcal{N}(0, I) \\&= \sqrt{1 - \beta} (\sqrt{1 - \beta} \mathbf{x}_{t-2} + \sqrt{\beta} \mathcal{N}(0, I)) + \sqrt{\beta} \mathcal{N}(0, I) \\&= (\sqrt{1 - \beta})^2 \mathbf{x}_{t-2} + \sqrt{1 - \beta} \sqrt{\beta} \mathcal{N}(0, I) + \sqrt{\beta} \mathcal{N}(0, I) \\&\dots \\&= (\sqrt{1 - \beta})^t \mathbf{x}_{t-t} + \dots + (\sqrt{1 - \beta})^2 \sqrt{\beta} \mathcal{N}(0, I) + \sqrt{1 - \beta} \sqrt{\beta} \mathcal{N}(0, I) \\&\quad + \sqrt{\beta} \mathcal{N}(0, I)\end{aligned}$$

Forward process: a closer look

$$\mathbf{x}_t = (\sqrt{1-\beta})^t \mathbf{x}_0 + \cdots + (\sqrt{1-\beta})^2 \sqrt{\beta} \mathcal{N}(0, I) + \sqrt{1-\beta} \sqrt{\beta} \mathcal{N}(0, I) + \sqrt{\beta} \mathcal{N}(0, I)$$

- for large t , $(\sqrt{1-\beta})^t$ approaches to 0
- each of $(\sqrt{1-\beta})^i \sqrt{\beta} \mathcal{N}(0, I)$ is a Gaussian with mean 0 and variance $(1-\beta)^i \beta$.
- All these Gaussian can be written as one Gaussian distribution with variance:

$$\sum_{i=0}^t (1-\beta)^i \beta = \beta \frac{1 - (1-\beta)^{t+1}}{1 - (1-\beta)} \sim \frac{\beta}{\beta} = 1$$

For large enough t we converge to the Normal distribution $\mathcal{N}(0, I)$.

Forward process: a closer look

- In practice we do not use a fixed noise β
- Linear schedule noise: $\beta_1 = 0.0001$ and $\beta_T = 0.02$
- The number of steps is about 10000 (application dependent), but this is slow!
- Recall $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \mathcal{N}(0, I)$.
- Let's define $\alpha_t = 1 - \beta_t$ and do the recursive expansion.

Forward process: a closer look

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \mathcal{N}(0, I) \\ &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \mathcal{N}(0, I) \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \mathcal{N}(0, I)) + \sqrt{1 - \alpha_t} \mathcal{N}(0, I) \\ &= (\sqrt{\alpha_t \alpha_{t-1}}) \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t + \alpha_t - \alpha_t \alpha_{t-1}} \sqrt{\beta_t} \mathcal{N}(0, I) \\ &\dots\end{aligned}$$

Forward process: a closer look

$$\begin{aligned}\mathbf{x}_t &= \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \mathcal{N}(0, I) \\&= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \mathcal{N}(0, I) \\&= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \mathcal{N}(0, I)) + \sqrt{1 - \alpha_t} \mathcal{N}(0, I) \\&= (\sqrt{\alpha_t \alpha_{t-1}}) \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t + \alpha_t - \alpha_t \alpha_{t-1}} \mathcal{N}(0, I) \\&\dots \\&= \sqrt{\alpha_t \alpha_{t-1} \cdots \alpha_2 \alpha_1} \mathbf{x}_0 + \sqrt{1 - \alpha_t \alpha_{t-1} \cdots \alpha_2 \alpha_1} \mathcal{N}(0, I) \\&= \sqrt{\prod_{i=1}^t \alpha_i} \mathbf{x}_0 + \sqrt{1 - \prod_{i=1}^t \alpha_i} \mathcal{N}(0, I) \\&= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \mathcal{N}(0, I) \\&= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon\end{aligned}$$

Reverse process: a closer look

Computing the reverse path (reverse denoising distribution) is not easy, but we know it is diffusion process.

We train a model to approximate the reverse distribution.

Similar to VAEs, the goal is to maximize a lower bound for the likelihood:

$$\log p(\mathbf{x}_0) \geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right]$$

A lot of effort goes into simplify this lower bound and exploiting the fact that we are dealing with diffusion process.

Reverse process: a closer look

At the end of the day, the model should learn the noise added

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{\mathbf{x}_0, t, \epsilon} [\|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|^2]$$

What does it mean? How do we actually do training ?

Reverse process: training

- Sample an image $\mathbf{x}_0$ and a time step t
- Sample a random noise ϵ .
- Get the noise image at time t : $\mathbf{x}_t = \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$
- Feed $\mathbf{x}_t$ to the neural network.
- The model's predicted noise ϵ_θ should be close to ϵ .

Algorithm 1 Training

- 1: **repeat**
 - 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
$$\nabla_\theta \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t) \right\|^2$$
 - 6: **until** converged
-

Reverse process: image generation

- Start from a noise image $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.
- Feed this to the trained neural network to predict a noise $\epsilon_\theta(\mathbf{x}_T)$
- Given this predicted noise we can do denoising and obtain $\mathbf{x}_{T-1}$ (we skipped through the math here)
- Repeat the above until we get to $\mathbf{x}_0$.

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

Semantic embeddings + diffusion

Text to image synthesis: Dall-E, Imagen, Stable Diffusion



“A chair that looks like a pineapple”

Transformers

We have seen semantic embeddings for words.

- An AI model.
- A fashion model.

The restaurant refused to serve me a ham sandwich because it only cooks vegetarian food. In the end, they just gave me two slices of bread. Their ambiance was just as good as the food and service.

- How to process this text into representation for down stream tasks ?
 - Positive or negative review ?
 - "Does the restaurant serve steak?"
- What are some challenges we face?

What are some challenges we face?

- Input can be large: 37 words each with 1024 dimensional representation gives $37 \times 1024 = 37888$. Fully connected NN is not feasible for long text.
- Different length text, hence different length representation.

Observation: The network should share parameters across different words in different positions; similar to convolution.

We need a model where parameters don't increase with input length.

What are some challenges we face?

The restaurant refused to serve me a ham sandwich because it only cooks vegetarian food. In the end, they just gave me two slices of bread. Their ambiance was just as good as the food and service.

- The word *their* must **attend to** the word *restaurant*.

Observation: There must be connections between the words.

The strength of these connections will depend on the words themselves.

The connections should span across the text.

Dot product self attention

Dot product self attention

Transformer acquires these properties using dot product self attention.

- Takes N inputs $\mathbf{x}_1, \dots, \mathbf{x}_N$ of size $D \times 1$ and returns N outputs of size $D \times 1$
- First, computes N values (no ReLU), one value for each input

$$\mathbf{v}_i = \mathbf{b}_v + \mathbf{W}_v \mathbf{x}_i \quad \text{with } \mathbf{b}_v \in \mathbb{R}^D, \quad \mathbf{W}_v \in \mathbb{R}^{D \times D}$$

Dot product self attention

Transformer acquires these properties using dot product self attention.

- Takes N inputs $\mathbf{x}_1, \dots, \mathbf{x}_N$ of size $D \times 1$ and returns N outputs of size $D \times 1$
- First, computes N values (no ReLU), one value for each input

$$\mathbf{v}_i = \mathbf{b}_v + \mathbf{W}_v \mathbf{x}_i \quad \text{with } \mathbf{b}_v \in \mathbb{R}^D, \quad \mathbf{W} \in \mathbb{R}^{D \times D}$$

- The i -th output, $\mathbf{sa}_i[\mathbf{x}_1, \dots, \mathbf{x}_N]$ is a *weighted sum* of over all values $\mathbf{v}_1, \dots, \mathbf{v}_N$.

$$\mathbf{sa}_i[\mathbf{x}_1, \dots, \mathbf{x}_N] = \sum_{m=1}^N a[\mathbf{x}_m, \mathbf{x}_i] \cdot \mathbf{v}_m$$

$a[\mathbf{x}_m, \mathbf{x}_i]$ is the *attention* the i -th output pays to input m .

Dot product self attention




Image from *Understanding Deep Learning* by Simon Prince

Attention weights

- Attention weights are non-linear functions of input.
- Compute N queries and N keys from input

$$\mathbf{q}_i = \mathbf{b}_q + \mathbf{W}_q \mathbf{x}_i$$

$$\mathbf{k}_i = \mathbf{b}_k + \mathbf{W}_k \mathbf{x}_i$$

- Compute similarities and pass through softmax:

$$a[\mathbf{x}_m, \mathbf{x}_i] = \text{softmax}_m(\text{sim}(\mathbf{k}_m, \mathbf{q}_i))$$

We can use vector dot product as a measure of similarity.

Attention weights

- Attention weights are non-linear functions of input.
- Compute N queries and N keys from input

$$\mathbf{q}_i = \mathbf{b}_q + \mathbf{W}_q \mathbf{x}_i$$

$$\mathbf{k}_i = \mathbf{b}_k + \mathbf{W}_k \mathbf{x}_i$$


- Compute similarities and pass through softmax:

$$\begin{aligned} a[\mathbf{x}_m, \mathbf{x}_i] &= \text{softmax}_m(\mathbf{k}_m^\top \mathbf{q}_i) \\ &= \frac{\exp(\mathbf{k}_m^\top \mathbf{q}_i)}{\sum_{j=1}^N \exp(\mathbf{k}_j^\top \mathbf{q}_i)} \end{aligned}$$

Attention weights

- For each input $\mathbf{x}_i$ the attention weights are positive and sum to one.
- The attention weights depend on the relative similarities between the i -th query and all of the keys.
- The queries and keys must have the same dimensions. However, these can differ from the dimension of the values, which is usually the same size as the input.
- Note there is no activation function, but the mechanism is nonlinear due to the dot-product and a softmax operation used to compute the attention weights.

Attention weights



What do we have so far?

- First, there is a single shared set of parameters $\phi = \{\mathbf{b}_v, \mathbf{W}_v, \mathbf{b}_q, \mathbf{W}_q, \mathbf{b}_k, \mathbf{W}_k\}$. This is independent of the number of inputs N , so the network can be applied to different sequence lengths.
- Second, there are connections between the inputs (words), and the strength of these connections depends on the inputs themselves via the attention weights.

Matrix form

- Data matrix $\mathbf{X}$
- $\mathbf{V} = \mathbf{b}_v \mathbf{1}^\top + \mathbf{W}_v \mathbf{X}$ ($\mathbf{1}$ is $N \times 1$)
- $\mathbf{Q} = \mathbf{b}_q \mathbf{1}^\top + \mathbf{W}_q \mathbf{X}$
- $\mathbf{K} = \mathbf{b}_k \mathbf{1}^\top + \mathbf{W}_k \mathbf{X}$
- $\mathbf{Sa}[\mathbf{X}] = \mathbf{V} \cdot \text{softmax}(\mathbf{K}^\top \mathbf{Q})$ performs the softmax operation independently on each of its columns
- Often we use the scaled version $\mathbf{Sa}[\mathbf{X}] = \mathbf{V} \cdot \text{softmax}\left(\frac{\mathbf{K}^\top \mathbf{Q}}{\sqrt{D_q}}\right)$.
Scaled attention converges faster and yields better results.

Matrix form




Image from *Understanding Deep Learning* by Simon Prince

Positional encoding

So far we have ignored the order of input and our set up is permutation equivariant.

The sentence **The woman ate the raccoon** has a different meaning than **The raccoon ate the woman**.

Two main approaches:

- Absolute positional encodings
- Relative positional encodings

$$\text{Sa}[\mathbf{X}] = (\mathbf{V} + \mathbf{\Pi}) \cdot \text{softmax}((\mathbf{K}^\top + \mathbf{\Pi})(\mathbf{Q} + \mathbf{\Pi}))$$

Multiple heads

Multiple self-attention mechanisms are usually applied in parallel, and this is known as multi-head self-attention. Now H different sets of values, keys, and queries are computed:

- Data matrix $\mathbf{X}$
- $\mathbf{V}_h = \mathbf{b}_{vh}\mathbf{1}^\top + \mathbf{W}_{vh}\mathbf{X}$ ($\mathbf{1}$ is $N \times 1$)
- $\mathbf{Q}_h = \mathbf{b}_{qh}\mathbf{1}^\top + \mathbf{W}_{qh}\mathbf{X}$
- $\mathbf{K}_h = \mathbf{b}_{kh}\mathbf{1}^\top + \mathbf{W}_{kh}\mathbf{X}$
- The h -th self-attention mechanism or head can be written as:
 $\mathbf{Sa}_h[\mathbf{X}] = \mathbf{V}_h \cdot \text{softmax}(\mathbf{K}_h^\top \mathbf{Q}_h)$

Multiple heads

Typically, if the dimension of the inputs $\mathbf{x}_i$ is D and there are H heads, the values, queries, and keys will all be of size D/H , as this allows for an efficient implementation. The outputs of these self-attention mechanisms are vertically concatenated, and another linear transform $\mathbf{W}_c$ is applied to combine them.

$$\mathbf{MhSa}[\mathbf{X}] =$$

$$\mathbf{W}_c[\mathbf{Sa}_1[\mathbf{X}]^\top, \mathbf{Sa}_2[\mathbf{X}]^\top, \dots, \mathbf{Sa}_H[\mathbf{X}]^\top]^\top$$




Image: *Understanding Deep Learning* by S. Prince

Transformer layers

In a real network, the data passes through a series of these transformer layers.




Image: *Understanding Deep Learning* by S. Prince

Transformers for natural language processing

A generic pipeline in NLP that uses transformers is as follows:

- **Tokenization:** split the text into common words and word fragments from which larger and less frequent words can be composed. Typical vocabulary size: 30,000
- **Embedding:** applies embedding to a bag-of-words matrix to get embedding for each word/token (previous lecture). Typical embedding dimension = 1024
- **Embedding** the embedding matrix goes through K transformers layer.
 - Encoder: transforms the text embeddings into a representation that can support a variety of tasks. Ex: BERT
 - Decoder: predicts the next token to continue the input text. Ex: GPT

Encoder model example: BERT

Architecture: A 24-layer Transformer with 16-head self-attention (each head uses 64-dim queries/keys/values. the matrices $\mathbf{W}_{vh}$, $\mathbf{W}_{qh}$, $\mathbf{W}_{kh}$ are 64×1024). The MLP size in each layer is 4096. Total of 340M parameters.

BERT pre-training

- Pre-training uses self-supervision, enabling training on massive text corpus without manual labels.
- BERT's self-supervised task: predicting masked (missing) words from internet text.
- Training setup:
 - Maximum input length: 512 tokens
 - Batch size: 256
 - 1 million training steps (about 50 epochs over a 3.3-billion-word corpus)
- Predicting missing words encourages the model to learn:
 - Syntax (e.g., “red” appears before nouns, not verbs)
 - Basic commonsense patterns (e.g., “train” is more likely than “peanut” in “The <mask> pulled into the station”)

Fine tuning

Model parameters are adjusted to specialize the pretrained network for a specific task. An additional output layer is added to map transformer outputs to the required prediction format.

Examples of fine-tuning tasks:

- **Text classification (e.g., sentiment analysis):**
 - A special <CLS> token is prepended during pre-training.
 - Its final hidden vector is mapped to a single value and passed through a sigmoid.
 - Trained using binary cross-entropy loss.
- **Word classification (e.g., named entity recognition):**
 - Each token embedding x_n is mapped to an $E \times 1$ vector (one score per entity type).
 - A softmax over the E classes gives per-token class probabilities.
 - Trained using multiclass cross-entropy loss.

Decoder model example: GPT3

- GPT-3 is an example of a **decoder-only** transformer model.
- Its architecture is similar to the encoder: stacked Transformer layers applied to learned word embeddings.
- Difference in purpose:
 - Encoder models: build text representations for downstream NLP tasks.
 - Decoder models: generate the **next token** in a sequence.
- By repeatedly feeding its own output back as input, the decoder can generate coherent text passages.

Decoder model example: GPT3

GPT-3 is an autoregressive language model: it predicts each token based on all previous tokens.




Image by: S. Prince

Language modeling

The probability of the entire sentence is factored into conditional probabilities:

$$\begin{aligned} Pr(\text{It takes great courage to let yourself appear weak}) = \\ Pr(\text{It}) \times Pr(\text{takes}|\text{It}) \times Pr(\text{great}|\text{It takes}) \times Pr(\text{courage}|\text{It takes great}) \times \\ Pr(\text{to}|\text{It takes great courage}) \times Pr(\text{let}|\text{It takes great courage to}) \times \\ Pr(\text{yourself}|\text{It takes great courage to let}) \times \\ Pr(\text{appear}|\text{It takes great courage to let yourself}) \times \\ Pr(\text{weak}|\text{It takes great courage to let yourself appear}). \end{aligned} \quad (12.14)$$

Example: *Understanding Deep Learning* by S. Prince

By multiplying these conditional probabilities, the model implicitly defines the **joint probability** of the full token sequence:

$$Pr(t_1, t_2, \dots, t_N) = \prod_{n=1}^N Pr(t_n \mid t_1, \dots, t_{n-1}).$$

Next-token prediction and joint probability modeling are two views of the same underlying task.

Predicting all next words simultaneously?

A decoder is trained by maximizing the log-probability of the next token at every position. How can we prevent “cheating” i.e. blocking access to future words?


How can we predict all the next words simultaneously using attention mechanism ?

Masked attention

Fast and scalable training that maintains the autoregressive constraints required for language modeling.

- **Masked self-attention prevents cheating:** Each token can only attend to itself and earlier tokens, blocking access to future words.
- **Enables parallel next-token prediction:** With the mask in place, the model can predict the next word at every position in the sequence in a single forward pass.
- **Efficient training:** All tokens contribute a supervised next-token prediction, allowing the model to learn from every position simultaneously.
- **Autoregressive behavior preserved:** Even though computation is parallel, each prediction still follows the correct conditional structure $Pr(t_n \mid t_1, \dots, t_{n-1})$.


Masked attention



Example: by S. Prince

- Sequence lengths are 2048 tokens long
- Batch size is 3.2 million tokens.
- 96 transformer layers, each of which processes a word embedding of size 12288.
- 96 heads in the self-attention layers and the value, query, and key dimension is 128.
- 300 billion tokens
- 175 billion parameters

LLM becoming very large indeed



from: <https://thelowdown.momentum.asia/the-emergence-of-large-language-models-llms/>